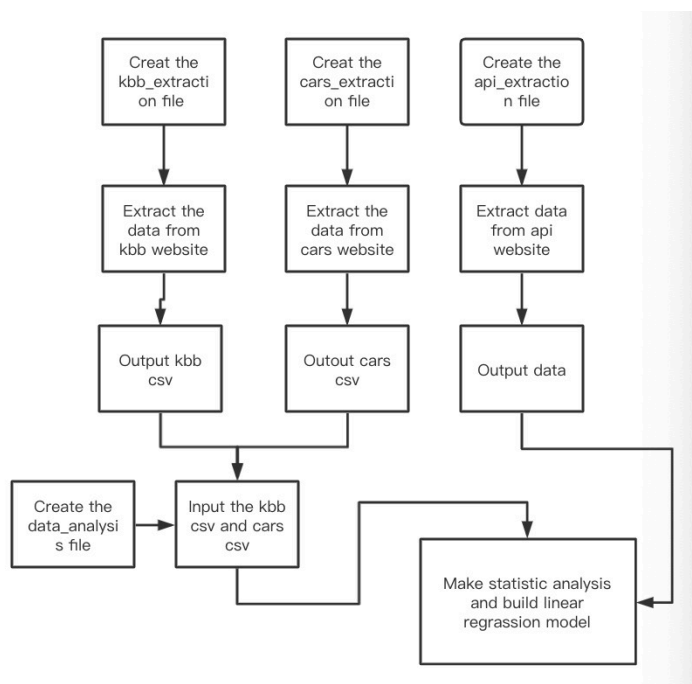


Final Project

Section 1 :

I want to get 1800 cars' price, mileage, year and other information from three well-known websites selling used cars in the United States. This project studies the influence of vehicle mileage on vehicle price and the influence of vehicle year on vehicle price, This project also studies the comprehensive influence of mileage and year on vehicle price. The linear regression algorithm is used to build the prediction model. I study the features of the most popular cars on the website, the color distribution of vehicles, and count the most and least colors on the website. This information can provide reference value for customers when they buy a car.

Section 2:



My code is divided into four files. the first file's name is 'kbb_extraction'. The main function of this file is to extract the data of 1000 Honda used cars. The output is a CSV file and the CSV file will be stored in the same place with the 'data_analysis' file. You can successfully import the data in the 'data_analysis' file. The second file's

name is 'cars_extraction'. The main function of this file is to extract the data of 800 BMW used cars. The output of a CSV file and the CSV file will be stored in the same place with the 'data_analysis' file. You can successfully import the data in the 'data_analysis' file. The third file's name is 'API_extraction'. It is to use the API to get the price of the new Honda series. The fourth file's

name is 'data_analysis'. The main function of this file is to import the two CSV files and make statistical analysis and establish a linear regression model.

Section 3:

a) kbb_extraction.py

For the kbb website, there are only 25 vehicles on each page, so I design a loop to process the URL. I change the page number displayed in the URL to the variable *i*. The variable *i* will increase every time in the loop, so the data of each page can be automatically extracted. When extracting the data of each page, I create five empty lists (*list_price*, *list_mile*, *list_color*, *list_model*, *list_year*) to store the data in order. And then, I connect to the website with the method of *requests.get*. I use the method of beautiful soup to get the HTML of this website and find that the data is in the class 'div '. Each type of data has a different class name, so I extract these data by looking up different class names and save these data in the empty lists I create. I process for the format of each type of data, and I change the digital data to int or float. Finally, I used a loop to combine each data in the five lists in order and extract it into a new list. The storage type is a large list containing many small lists. Finally, I used PD. Dataframe to convert these data into CSV file.

```

from bs4 import BeautifulSoup
import urllib
import requests
import pandas as pd

def extract():
    i=0
    data_list=[]
    while i<1000:
        a=1
        list_price=[]
        list_mileage=[]
        list_color=[]
        list_model=[]
        list_year=[]
        url=('https://www.kbb.com/cars-for-sale/used/honda/accord/los-angeles'+
            '-ca-90001?makeCodeList=HONDA&searchRadius=500&modelCodeList=ACCORD&'+
            'zip=90001&marketExtension=include&listingTypes=USED&isNewSearch=true'+
            '&sortBy=relevance&numRecords=25&firstRecord='+str(i))#url
        website=requests.get(url)#connect website
        soup=BeautifulSoup(website.content,'html.parser')
        tag1=soup('div')#find class
        for tag in tag1:#extract data
            if tag.get('class',None):
                if tag.get('class',None)[0]=='text-gray-base':
                    if 'MSRP' in tag.text.split(',')[1] and '$' not in tag.text.split(',')[1]:
                        list_price.append(float(tag.text.split(',')[0]+
                            tag.text.split(',')[1].rstrip('MSRP')))
                    elif '$' in tag.text.split(',')[1]:
                        list_price.append(float(tag.text.split(',')[0]+
                            tag.text.split(',')[1].split(' ')[0].rstrip('MSRP')))
                else:
                    list_price.append(float(tag.text.split(',')[0]+
                        tag.text.split(',')[1]))
            if tag.get('class',None)[0]=='text-bold':
                if 'miles' in tag.text:
                    if len(tag.text.split(' ')[0].lstrip('('))<5:
                        list_mileage.append(float(tag.text.split(' ')[0].lstrip('(')))
                    else:
                        list_mileage.append(float(tag.text.split(' ')[0].split(',')[0]
                            +tag.text.split(' ')[0].split(',')[1]))
                if tag.get('class',None)[0]=='display-flex':
                    if 'Accord' in tag.text:
                        list_year.append(int(tag.text.split(' ')[1]))
                        list_model.append(tag.text.split(' ')[3])
                if tag.get('class',None)[0]=='item-card-specifications':
                    list_color.append(tag.text.split(':')[1])
        while a<28:
            data_list.append([list_year[a],list_model[a],list_mileage[a-1],list_price[a],list_color[a]])
            #total of data
            a=a+1
        i=i+25
    print(data_list)
    name=['year','model','mileage','price','color']
    file=pd.DataFrame(columns=name,data=data_list)
    file.to_csv('kbb_data.csv',index=False,encoding="utf-8")#to csv file
extract()

```

b) cars_extraction.py

For cars website, there are 100 cars on each page, so I take the same method as above to extract the data of 800 cars.

```
from bs4 import BeautifulSoup
import urllib
import requests
import pandas as pd

def get_data():
    i=1
    data_list=[]
    while i<11:
        a=0
        list_price=[]
        list_mileage=[]
        list_color=[]
        list_model=[]
        list_review=[]
        list_year=[]
        website=requests.get('https://www.cars.com/for-sale/searchresults.action/?mdId=20444&mkId=20005&page='+str(i)
                            +'&perPage=100&rd=99999&searchSource=PAGINATION&sort=relevance&stkTypeId=28881&z=90007')
        soup=BeautifulSoup(website.content,'html.parser')
        tag1=soup('span')
        for tag in tag1:
            if tag.get('class',None):
                if tag.get('class',None)[0]=='listing-row__price':
                    list_price.append(float(tag.text.strip('\n ').rstrip('$').split(',')[0]
                                           +tag.text.strip('\n ').rstrip('$').split(',')[1]))
                if tag.get('class',None)[0]=='listing-row__mileage':
                    list_mileage.append(float(tag.text.strip('\n ').rstrip(' mi.').split(',')[0]
                                           +tag.text.strip('\n ').rstrip(' mi.').split(',')[1]))
                if tag.get('class',None)[0]=='listing-row__review-number':
                    list_review.append(int(tag.text.strip('\n ').strip(' ').rstrip(' reviews')))
        tag2=soup('h2')
        for tag in tag2:
            if tag.get('class',None):
                if tag.get('class',None)[0]=='listing-row__title':
                    year=tag.text.strip('\n ').split(' ')[0]
```

```

        model=(tag.text.strip('\n ').split(' ')[1]+tag.text.strip('\n ').split(' ')[2]
                +tag.text.strip('\n ').split(' ')[3])
        list_year.append(int(year))
        list_model.append(model)
    tag3=soup('ul')
    for tag in tag3:
        if tag.get('class',None):
            if tag.get('class',None)[0]=='listing-row__meta':
                list_color.append(tag.text[55:62].strip('\n '))
    while a<80:
        data_list.append([list_year[a],list_model[a],list_mileage[a],
                           list_review[a],list_color[a],list_price[a]])
        a=a+1
    i=i+1
    print(data_list)
    name=['year', 'model', 'mileage', 'review', 'color', 'price']
    file=pd.DataFrame(columns=name,data=data_list)
    file.to_csv('cars_data.csv',index=False,encoding="utf-8")
get_data()

```

c) api_extraction.py

```

import json
from bs4 import BeautifulSoup
import urllib
import requests
import pandas as pd
def get_cars_infomation(make,model):
    data_list=[]
    api_key='Z7sq4yf1s4IjVLgzZfVhWU2LiKdVHScy'#get api key
    j=2020
    while j>2019:# extrat 2020 data
        list_price=[]
        list_mileage=[]
        list_model=[]
        list_year=[]
        a=1
        url=('http://marketcheck-prod.apigee.net/v2/search/car/active?api_key='+api_key
              +'&year='+str(j)+'&make='+make +'&model'+model)
        resp=requests.get(url)#connect url
        js=resp.json()# covert json formal
        prettyjs=json.dumps(js, indent=4, sort_keys=True)
        for i in range(10):
            list_model.append(js['listings'][i]['build']['model'])#model data
            list_year.append(js['listings'][i]['build']['year'])#year data
            list_price.append(js['listings'][i]['price'])#price data
            list_mileage.append(js['listings'][i]['ref_miles'])#mileage data
        while a<9:
            data_list.append([list_year[a],list_model[a],list_mileage[a],list_price[a]])# total of data
            a=a+1
        j=j-1
    name=['year', 'model', 'mileage', 'price']
    return data_list
get_cars_infomation('honda','accord')

```

Firstly, I apply for a free API key from the API website. Then I build four empty lists to store the data(list_price, list_mileage, list_model, list_year). I use the method of requests.get to connect with websites. And then I get the html and convert it into JSON format to help my data search. According to the different positions of each data, I take them out by using the method of the dictionary and store them in four lists. Finally, I used a loop to combine each data in the four lists in order and extract it into a new list. The storage type is a large list with many small lists.

d) data_analysis.py

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm

def kbb_analysis():
    print('The relationship between mileage and price:')
    data=pd.read_csv('kbb_data.csv',header=0)#import kbb website data
    X1 = data['mileage'].values.reshape(-1,1)
    y1 = data['price'].values.reshape(-1,1)
    X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=0)
    #split data
    reg = LinearRegression()
    # linear regression
    reg.fit(X1_train, y1_train)
    print('coefficient:',reg.coef_, 'intercept:',reg.intercept_)
    #print coefficient and intercept
    y1_pred = reg.predict(X1_test)
    #predict y1
    plt.scatter(X1_test, y1_test, color='gray')
    plt.plot(X1_test, y1_pred, color='red', linewidth=2)
    plt.xlabel("mileage")
    plt.ylabel("price")
    plt.show()
    print('The relationship between year and price:')
    X2=data['year'].values.reshape(-1,1)
    y2 = data['price'].values.reshape(-1,1)
    X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.2, random_state=0)
    #split data
    reg.fit(X2_train, y2_train)
    y2_pred = reg.predict(X2_test)
    #predict y2
```



```

print('coefficient:', reg.coef_, 'intercept:', reg.intercept_)
#print coefficient and intercept
plt.scatter(X2_test, y2_test, color='gray')
plt.plot(X2_test, y2_pred, color='red', linewidth=2)
plt.xlabel("year")
plt.ylabel("price")
plt.show()
print('The relationship between mileage, year and price:')
X=data[['mileage', 'year']]
y=data['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
#split data
reg.fit(X_train, y_train)
y_pred = reg.predict(X_test)
#predict y
df = pd.DataFrame({'Actual price': y_test, 'Predicted price': y_pred})# predict price
print(df)
df_25 = df.head(25)#get 25 data
df_25.plot(kind='bar', figsize=(10,8))
plt.grid(which='major', linestyle='-', linewidth='0.5', color='green')
plt.grid(which='minor', linestyle=':', linewidth='0.5', color='black')
plt.show()
model = LinearRegression().fit(X, y)
model = sm.OLS(y, X)#linear regression
results = model.fit()
print(results.summary())#parameters of linear regression
kbb_analysis()

```

Firstly, the data in CSV format is imported into python by read_CSV method. The relationship between vehicle mileage and price should be studied. So the mileage data is given to X1 and the price data is given to y1. I use the method of train_test_split to divide the data into training data and test data. And I use the package of linear regression of sk-learn to get the function slope, intercept and drawing. Secondly, I study the relationship between vehicle year and price. So I give the mileage data to X2 and the price data to y2. I use the method of train_test_split to divide the data into training data and test data. I use the package of the linear regression of sk-learn to get the function slope, intercept and drawing. Thirdly, I study the comprehensive relationship of vehicle year, mileage and price. So I give the mileage and year data to X and the price data to y. I use the method of train_test_split to divide the data into training data and test data. I use sklearn's linear regression package so, I can use the predict function to generate the forecast price and draw a picture. The parameters form of linear regression are generated by the package of OLS.

Section4:

In the data_analysis file, I use two packages, linear regression package from sk-learn, which we don't learn in class. It can perform linear regression on the data. The other one is statsmodels.api, which is used to get all the parameters of linear regression.

Section 5:

step 1: `cd Desktop/final_6475336873`

step 2: `pip install -r requirements.txt`

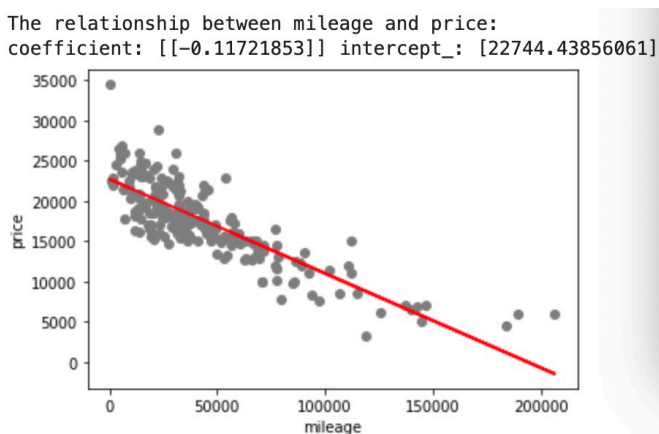
step 3: `cd code`

step 4: `input python3 kbb_extraction.py , python3 cars_extraction.py, python3 api_extraction.py , python3 data_analysis.py`

Section 6:


For Honda Accord:

From this graph, it can be seen that the mileage is inversely proportional to the price through univariate linear regression.




From this graph, we can see that the year is directly proportional to the price through univariate linear regression.

The relationship between year and price:
 coefficient: `[[1572.09370805]]` intercept_: `[-3152179.3677001]`



Through multiple linear regression, 25 second-hand cars are randomly selected from the sample to predict their prices. From this figure, it can be seen that the price predicted by using the model is probably accurate with little error.

[216 rows x 2 columns]



I use api to search the new Honda Accord' price is 26575. Therefore, no cars exceed 25000

in the picture.

```
: [[2020, 'Civic', 0, 20078],  
  [2020, 'Insight', 0, 25265],  
  [2020, 'Civic', 0, 20196],  
  [2020, 'Odyssey', 10, 35785],  
  [2020, 'Odyssey', 0, 37800],  
  [2020, 'Civic', 10, 23180],  
  [2020, 'Accord', 0, 26575],  
  [2020, 'Civic', 0, 24115]]
```


Through the parameter analysis of multiple linear regression, the square of R is 0.978, which shows that the model fits well.

OLS Regression Results						
Dep. Variable:	price	R-squared (uncentered):	0.978			
Model:	OLS	Adj. R-squared (uncentered):	0.978			
Method:	Least Squares	F-statistic:	2.413e+04			
Date:	Tue, 12 May 2020	Prob (F-statistic):	0.00			
Time:	19:13:23	Log-Likelihood:	-10081.			
No. Observations:	1080	AIC:	2.017e+04			
Df Residuals:	1078	BIC:	2.018e+04			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
mileage	-0.1169	0.002	-47.037	0.000	-0.122	-0.112
year	11.2724	0.066	171.984	0.000	11.144	11.401
Omnibus:	130.585	Durbin-Watson:	1.738			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	231.477			
Skew:	0.778	Prob(JB):	5.44e-51			
Kurtosis:	4.650	Cond. No.	41.8			

For BMW328i:


From this graph, we can see that the mileage is inversely proportional to the price.

The relationship between mileage and price:
coefficient: $[-0.12095285]$ intercept_: $[21289.001637]$




From this graph, we can see that the year is directly proportional to the price through univariate linear regression.

coefficient: $[1448.14127061]$ intercept_: $[-2901866.41057718]$



Through multiple linear regression, 25 second-hand cars are randomly selected from the sample to predict their prices. From this figure, it can be seen that the price predicted by using the model is probably accurate with little error.

[160 rows x 2 columns]



Through the parameter analysis of multiple linear regression, the square of R is 0.962, which shows that the model fits well.

OLS Regression Results

Dep. Variable:	price	R-squared (uncentered):	0.962
Model:	OLS	Adj. R-squared (uncentered):	0.962
Method:	Least Squares	F-statistic:	1.014e+04
Date:	Tue, 12 May 2020	Prob (F-statistic):	0.00
Time:	19:05:57	Log-Likelihood:	-7464.1
No. Observations:	800	AIC:	1.493e+04
Df Residuals:	798	BIC:	1.494e+04
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
mileage	-0.1206	0.003	-38.378	0.000	-0.127	-0.114
year	10.5822	0.115	92.129	0.000	10.357	10.808

Omnibus:	20.695	Durbin-Watson:	1.930
Prob(Omnibus):	0.000	Jarque-Bera (JB):	29.731
Skew:	0.246	Prob(JB):	3.50e-07
Kurtosis:	3.806	Cond. No.	87.6

By selecting the 25 cars with the most of reviews and calculating their means, I find that BMW 328 car with the year of 2013, the mileage of 65420 and the price of 14170 is very popular on the website.

	year	model	mileage	review	color	price
148	2013	BMW328i	27349.0	21024	White	15900.0
216	2011	BMW328i	69721.0	20415	Gray	9995.0
256	2016	BMW328328i	34649.0	15983	Silver	17729.0
168	2016	BMW328i	42788.0	8748	White	16991.0
156	2013	BMW328i	63109.0	4964	Brown	10989.0
315	2011	BMW328i	80927.0	4670	Black	9500.0
274	2008	BMW328xi	70163.0	4099	Black	7500.0
791	2016	BMW328i	41994.0	3763	Gray	25494.0
431	2008	BMW328xi	90057.0	3610	Gray	7488.0
690	2013	BMW328i	81703.0	3556	Black	12995.0
143	2016	BMW328i	41396.0	3413	White	14994.0
731	2012	BMW328i	108570.0	3165	White	10400.0
63	2016	BMW328i	32000.0	2613	Blue	18997.0
711	2014	BMW328i	68597.0	2474	Black	13250.0
41	2012	BMW328i	70535.0	2358	Blue	12500.0
365	2010	BMW328i	136494.0	2263	Blue	5500.0
347	2016	BMW328i	28786.0	2156	Gray	22450.0
348	2016	BMW328i	7399.0	2156	Silver	22995.0
739	2011	BMW328i	101942.0	2156	Black	8995.0
439	2016	BMW328i	15564.0	2152	Silver	22888.0
438	2011	BMW328i	72908.0	2152	Other	9499.0
519	2016	BMW328i	48018.0	2152	Blue	17450.0
364	2016	BMW328i	21088.0	2149	White	24880.0
414	2007	BMW328xi	169002.0	2148	Red	4995.0
516	2012	BMW328i	110742.0	2145	Black	9900.0
year	2013.04					
mileage	65420.04					
review	5059.36					
price	14170.96					
dtype:	float64					

The most popular colors of BMW 328i series on the website are black, white and gray. It is recommended not to buy green, gold and brown. They are very rare and hard to sell.

